

Olinethra Website — Design Brief for Interns

Oct 4, 2026 · @Olinethra

The idea

Olinethra is a marketing website for a software engineering and AI consultancy. It should feel like looking at a live production system, not a brochure. Every page borrows the look of developer tools: API endpoints as section labels, terminal windows, uptime numbers and status lights.

Who it is for: CTOs, VPs of Engineering and platform leads at mid-size and enterprise companies. They are technical, skeptical of sales talk, and want proof.

What a visitor should feel: "These people build serious systems, and they talk to me like an engineer."

What a visitor should do: book a consultation with an architect, or send a project brief. Every page ends by pushing toward one of those two actions.

Three ideas run through the whole design:


- **The site is the API.** Sections are labelled like endpoints (`GET /api/v1/services` , `POST /api/v1/inquiries`). The brand promise is shown through the interface itself.
- **Proof over adjectives.** Each claim carries a number: latency, uptime, throughput, cost saved.
- **No sales gatekeepers.** Copy repeats that you talk directly to a principal architect, never an account manager.

Site map

The site has four pages, and every path ends at Contact, where a visitor books a call or sends a brief.

Every page leads the visitor to the Contact page

Amber outline: Services exists, but no nav link reaches it yet



site map · 4 pages, 1 goal

Home sends visitors to Services and Projects, both of which close with a call to action that opens Contact. The orange header button is a shortcut to Contact from anywhere.

Page by page

Each page follows the same rhythm: a hero with a big claim, sections of proof, then a call to action. Sections are listed top to bottom in the order they appear.

The screenshots below show each page at desktop width (1280px). One limitation: the office photo on Contact did not load in the screenshot tool, so it shows as a grey block.

1. Home (`index.html` , same as `code.html`)

Job: win attention in 5 seconds and send the visitor to Services, Projects or a booking.

- Hero** — headline "Engineering Scalable Software & Intelligent Cloud Systems for Ambitious Brands", two buttons (Explore Services, View Case Studies) and three proof badges (SOC-2, <250ms P99, 40+ nodes).
- Live terminal** — a fake console with two tabs (`/health` , `/services`) showing JSON and a latency sparkline. This is the signature visual.
- Metrics strip** — four big numbers: 99.99% SLA, <250ms, 40+ deployments, SOC-2.
- Tech stack chips** — Next.js, TypeScript, Rust, Go, Python, Kubernetes.
- Four service cards** — Web & Next.js, Cloud & DevOps, Applied AI, APIs & Distributed Systems. Each has tags and one benchmark.
- Three case study cards** — FinScale, Aether Health, Nexus Logistics, each with two metrics and tech badges.
- Testimonials** — three five-star quotes from engineering leaders.

8. **Call to action** — "Book Principal Discovery" and "Request System Audit".
9. **AI chat drawer** — a floating assistant opened from the bottom-right button.

OLINETHRA

ENGINEERING & AI

API Operational v1

Services

Projects

Contact

Client Portal

Schedule Consultation

ENTERPRISE CLOUD ARCHITECTURE & AI RUNTIME

v2.4.0-stable

Engineering Scalable Software & Intelligent Cloud Systems for Ambitious Brands

We architect next-generation web platforms, enterprise AI solutions, and distributed systems with uncompromising speed, security, and precision. Built for 99.999% reliability under extreme load.

Explore Services →

View Case Studies

SOC-2 Type II Certified

<250ms Global P99

40+ Multi-Region Nodes

SERVICE GUARANTEE

99.99%

Guaranteed SLA Uptime Target

EDGE DISPATCH

<250ms

Global P99 Worldwide Edge

DEPLOYMENTS

40+

Fortune 1000 & High-Growth Scalars

COMPLIANCE

SOC-2

Type II & HIPAA Hardened

CORE INFRASTRUCTURE FRAMEWORKS

Battle-Tested Modern Stack

Designed with zero legacy dependencies and clean architecture principles.

N

Next.js 15

Turbopack SSR

TS

TypeScript

Strict Typing

Rs

Rust

Zero-Cost Mem

Go

Golang

gRPC Microservices

Py

Python / PyTorch

LLM Pipelines

K8s

Kubernetes

Multi-region mesh

GET /api/v1/services — CAPABILITY REGISTRY

Specialized Engineering Disciplines

Every architecture is built with fault tolerance, sub-second streaming, and enterprise governance baked into the core runtime.

Core Web Vitals 100/100

Full-Stack Web & Next.js Platforms

Ultra-responsive distributed web architectures built with React Server Components, zero-layout-shift design, and smart edge caching for frictionless consumer journeys.

Micro-frontends

Edge Rendering

Dynamic Hydration

Benchmark

<45ms First Contentful Paint

→

Multi-Cloud Failover

Enterprise Cloud & DevOps

Resilient containerized infrastructure engineered for continuous automated delivery. Immutable Terraform blueprints, zero-downtime rolling updates, and self-healing Kubernetes clusters.

Kubernetes EKS/GKE

Terraform IaC

GitOps ArgoCD

Target RTO/RPO

<60s Automated Recovery

→

Autonomous Intelligence

100k+ Events/sec

edge-us-east-1.olinethra.internal

STATUS: 200 OK

GET /api/v1/health

GET /api/v1/services

{

"status": "healthy",

"cluster_runtime": "k8s-v1.31.2-arm64",

"active_nodes": 48,

"p99_latency": "18.4ms",

"sla_adherence": "99.999%",

"llm_inference_load": "normal"

}

Realtime P99 Telemetry

18.42 ms

Page 4 of 21

Applied AI & LLM Systems

Domain-fine-tuned neural agents, hybrid vector retrieval-augmented generation (RAG), and strict safety boundary evaluation for compliant enterprise intelligence.

Fine-Tuned LLMs Vector DBs (Qdrant) /api/v1/chat Engine

Accuracy Metric
99.4% Hallucination Suppression



High-Throughput APIs & Distributed Systems

Decoupled event-driven backends, lock-free concurrency in Go and Rust, and ultra-compact binary gRPC protocols architected for massive concurrent transactions.

Kafka Streams GraphQL Federation Distributed Cache

Network Overhead
<4ms Serialization Overhead



PRODUCTION ARTIFACTS

Enterprise Deployments & Case Studies

[Request Full Engineering Casebook >](#)

FINTECH & HFT

FinScale Global

Sub-millisecond data feeds and algorithmic order processing handling \$4.2B in weekly equities volume with zero system downtime.

Throughput

+320%

Latency

0.8ms P99

Rust Next.js 15 WebSocket Redis Cluster

[Inspect Architecture Brief →](#)

HEALTHCARE AI

Aether Health

End-to-end encrypted telemedicine core handling 500k+ monthly consultations, automated clinical summarization, and zero-leakage records.

Latency Reduction

4.2x Faster

Uptime

99.999%

Python Go WebRTC AWS KMS

[Inspect Architecture Brief →](#)

SUPPLY CHAIN ML

Nexus Logistics

Predictive routing AI mesh evaluating 18M daily telemetry checkpoints to dynamically circumvent maritime and land route disruptions.

Transit Delay

-34% Less

Compute Cost

-48% AWS

Kafka PostgreSQL Kubernetes PyTorch

[Inspect Architecture Brief →](#)

ENGINEERING LEADERSHIP REVIEWS

Trusted by Technical Decision Makers

★★★★★

"Olinethra's distributed systems group shaved 80ms off our global P99 in our first production deployment. Their Rust and Kafka pipeline design is unmatched in rigor."

MR

Marcus Reynolds

VP of Infrastructure, Horizon Fin

★★★★★

"Their LLM RAG pipelines took our enterprise knowledge engine from an experimental prototype to 99.4% verifiable accuracy in heavily regulated clinical environments."

EV

Elena Vargas, Ph.D.

Chief Technology Officer, BioSynapse

★★★★★

"Zero-downtime multi-region Kubernetes migration carried out across three cloud providers without a single dropped packet. Outstanding technical execution."

SK

Siddharth Kumar

Head of Platform, OmniCore Logistics

📅 Q3 & Q4 Architecture Sprint Openings

Ready to Architect Your Next

📅 Book Principal Discovery

Technical Breakthrough?

Book a 45-minute technical discovery session directly with a Principal Cloud Systems Architect. No sales intermediaries — pure architecture, specs, and timeline estimation.

[Request System Audit](#)



OLINETHRA

High-performance distributed systems, low-latency microservices, and specialized autonomous AI agents engineered for enterprise-grade workloads.

Core API Gateway: 99.99% Uptime

SOLUTIONS

Autonomous Agents
Distributed Pipelines
Neural Search Engines
Edge AI Deployment
Security Auditing

ENGINEERING BRIEFING

Receive our bi-weekly architectural dispatches covering production AI workloads, memory optimizations, and system design.

engineer@domain.com

Subscribe

© 2025 Olinethra Technologies Inc. All systems verified and reserved.

[Privacy Protocol](#) [Service SLA](#) [API Telemetry](#)

2. Services (`code3.html`)

Job: explain what we do in depth and help a buyer size a project.

1. **Header** — "End-to-End Engineering & Modern Software Capabilities" with a cluster status card (14 regions, 99.995%).
2. **Four pillar cards with filter tabs** — Frontend, Cloud, AI, SRE. Each card lists four capabilities and three deliverables.
3. **Production flow diagram** — Edge Next.js → Go core → RAG → Zero-Trust SLA, plus four feature tiles.
4. **Delivery lifecycle** — four phases: Discovery (sprints 1-2), Blueprint (sprint 3), Bi-weekly shipments (sprints 4-12), Launch and SLA.
5. **Scope estimator** — pick architecture type, sprint count and security tier; the card on the right updates team, launch sprint and SLA.
6. **Call to action banner.**



OLINETHRA
ENGINEERING &
AI

API Operational
v1

Services

Projects

Team

Insights

Careers

FAQs

Contact

Client
Portal

Schedule
Consultation



GET /API/V1/SERVICES — READY

Latency: 14ms / Coverage: 100% Zero-Downtime

End-to-End Engineering & Modern Software Capabilities

From architectural blueprint to global production scale, we design, build, and deploy resilient digital software products engineered for zero compromise.

Active Production Clusters14 Region Nodes

Availability: 99.995%SOC-2 Type II Validated

ARCHITECTURE DISCIPLINES

Precision Pillars & Core Systems

All DisciplinesNext.jsCloud SystemsApplied AISRE & DevOps



SYS.01 // WEB_RUNTIME

Next.js & Frontend Architecture

Ultra-responsive enterprise interfaces built on React Server Components, zero-waterfall data fetching pipelines, and automated sub-second Core Web Vitals guarantees.

CORE TECHNICAL CAPABILITIES

Server Actions & React 19 RSC

Micro-frontend dynamic federation

Core Web Vitals < 0.8s LCP

WCAG 2.1 AAA accessible systems

ENGINEERED DELIVERABLES

Production PWA / SPA

Design Token Library

Automated Playwright E2E



SYS.02 // DISTRIB_SVC

Custom Cloud & Distributed Backend Systems

High-concurrency microservices, multi-region distributed databases, and event streams engineered for millions of requests per second with deterministic latency.

CORE TECHNICAL CAPABILITIES

Kubernetes Mesh Orchestration

Kafka & Serverless Event Buses

High-Concurrency Go & Rust daemons

Postgres / Citus global sharding

ENGINEERED DELIVERABLES

Multi-Region Cloud Infra

GitOps CI/CD Pipelines

Datadog & Grafana Dashboards



SYS.03 // COGNITIVE_AI

Applied Artificial Intelligence & Machine Learning

Domain-specific neural pipelines, localized retrieval-augmented generation (RAG), and autonomous multi-agent systems built directly into core enterprise workflows.

CORE TECHNICAL CAPABILITIES

Domain LLM Fine-Tuning (LoRA)

Vector Stores (Pinecone / Milvus)

Hybrid BM25 + Dense RAG Search

LangGraph Autonomous Orchestrators

ENGINEERED DELIVERABLES

Secure Private LLM Gateway

Embeddings Ingestion Pipeline

Custom Evaluation & Guardrails



SYS.04 // RESILIENCE

DevOps, Security & Site Reliability Engineering

Hardened infrastructure posture adhering to Zero-Trust principles, automated disaster recovery simulations, and 24/7 proactive incident response capabilities.

CORE TECHNICAL CAPABILITIES

Zero-Trust IAM & Network Policies

SOC-2 Type II Compliance Readiness

Chaos Engineering Experiments

Automated Pen-Testing in CI

ENGINEERED DELIVERABLES

Disaster Recovery Runbook

Automated Rollback Engine


Terraform & OpenTofu Modules

TOPOLOGY & PIPELINE EXECUTION

Integrated Production Flow

How our four core engineering specializations unify inside an automated delivery perimeter.


Page 7 of 21



METHODOLOGY & CADENCE

Agile Architecture & Delivery Lifecycle

Predictable, sprint-based release cycles designed to derisk novel technology early and sustain velocity to production handoff.



INTERACTIVE RESOURCE PLANNER

Interactive Pricing & Scope Estimator

Configure your technical scope requirements to generate dynamic sprint estimates and architecture team compositions.

Estimator Engine v2.4 Live

1. SYSTEM ARCHITECTURE TYPE

Frontend / Next.js

Modern RSC app with token design system.

Fullstack Cloud

Frontend + Distributed Go/Node backend.

AI Core & Microservices

LLM, vector stores, autonomous agent loops.

2. TARGET DELIVERY HORIZON

Rapid Sprint Spike (4 Sprints)

Balanced (8-10 Sprints)

Enterprise Platform (16 Sprints)

Standard Velocity (8 Sprints)

3. SECURITY POSTURE & COMPLIANCE REQUIREMENT

Standard Cloud

OWASP top-10 & TLS 1.3

SOC-2 Type II

Audit trails & zero-trust policies

HIPAA / FinTech

HSM encryption & isolated VPCs

Scope Telemetry

Estimated Range

TOTAL SPRINT COMMITMENT

8 Two-Week Sprints (~16 Weeks)

Pod: 1 Principal, 2 Staff Engineers, 1 DevOps

Architecture RFC Sign-off:

Sprint 1 - 2

Production Launch Target:

Sprint 8 Completion

Guaranteed Uptime SLA:


99.95% Availability

Security Assurance:

SOC-2 Type II Blueprint

Inquire With This Scope →

Estimates do not constitute binding quotes. Exact pods calibrated after Discovery Spike.



3. Projects (`code2.html`)

Job: prove we have done this before.

1. **Header** with three summary metrics (148 microservices, 14.2ms, \$1.8B+).
2. **Filter pills and live search** across six categories.
3. **Six case study cards** — FinScale Treasury, Aether Telehealth, OmniRoute, Synthetix Data Cloud, PulseVibe, Novus Retail. Each card shows client, challenge, two KPIs and stack.
4. **Empty state** when no card matches.
5. **Cumulative metrics** — availability, response time, cloud savings, zero critical defects.
6. **Developer API banner** with a sample `curl` call.
7. **Architecture modal** — "Inspect Architecture" opens a detail panel styled like an API response.



● API Operational v1

Services

Projects

Contact

Client Portal

Schedule Consultation



● PRODUCTION DEPLOYMENTS & ARCHITECTURE

Engineered for Impact: Case Studies & Production Systems

Explore real-world software architecture, digital transformations, and high-performance applications delivered by Olinethra for mission-critical enterprise workloads.

ACTIVE MICROSERVICES
148

AVG LATENCY SLA
14.2ms

TRANSACTIONS PROCESSED
\$1.8B+

All Projects (12)

Cloud & Infrastructure

AI & Machine Learning

Fintech & Web3

Healthcare Tech

Ent

 Search stack, client, or architecture

Production Case Studies

GET /api/v1/projects

● Latency: 28ms



Fintech & Web3

GO-LEDGER-V3

FinScale Global Bank

ISO 20022 Compliant

FinScale Treasury Platform

Real-time settlement and automated double-entry ledger powered by distributed event-...

 The Architecture Challenge

Eliminate 4hr legacy batch sync & eliminate split-brain write conflicts under high volume.

Settlement Volume
\$42M/day

Latency Floor
12ms avg

GoKafkaClickHouseDockerFastAPI

Inspect Architecture



Healthcare Tech

E2E-ENCRYPTED

Aether MedSystems

WebRTC Mesh

Aether Telehealth Core

WebRTC encrypted video pipelines and automated patient triaging with Next.js,...

 The Architecture Challenge

Sub-100ms video roundtrips across intermittent rural 4G connections without packet leakage.

Consultations
180k/mo

Intake Wait
-55%

TypeScriptNext.jsWebSocketsFastAPI

Inspect Architecture



AI & Machine Learning

GRAPH-RL-V4

OmniRoute Freight

Ray Clusters

OmniRoute AI Logistics

Autonomous routing engine processing 2M daily delivery coordinates with multi-agent...

 The Architecture Challenge

Solve NP-hard Traveling Salesperson Variations dynamically in sub-10ms intervals.

Delivery Coordinates
2.4M/day

OpEx Reductions
-45%

PythonPyTorchReactDocker

Inspect Architecture



Cloud & Infrastructure

TEE-ENCLAVE

Synthetix Labs

Confidential Compute

Synthetix Data Cloud

Zero-trust enterprise data cleanroom enabling secure multi-party analytics with cryptographic...

 The Architecture Challenge

Perform heavy join aggregations on encrypted



Enterprise SaaS

SOCKET-MESH

PulseVibe Media Inc.

Global CDN

PulseVibe Creator Economy

Global low-latency payout and streaming distribution network processing micro-...

 The Architecture Challenge

Handle bursty fan traffic spiking from 1k to



Enterprise SaaS

EDGE-RENDER

Novus Global Retail

99/100 Lighthouse


Novus Retail Engine

Headless omnichannel commerce delivering 99/100 Lighthouse performance with edge-...

 The Architecture Challenge

Support 50,000 requests/sec flash sale

Page 10 of 21



4. Contact (code4.html)

Job: turn interest into a lead.

1. Header — "Let's Build Something Exceptional Together".

2. **Left column** — WhatsApp direct line, uptime card, SLA promise (NDA in 24 hours), two email inboxes, three office hubs with image.
3. **Right column** — the project brief form: name, email, company, website, scope, budget, timeline, specs, file upload.
4. **Pre-brief AI agent** — an inline chat demo with suggested questions.

OLINETHRA

ENGINEERING & AI

API Operational v1

Services

Projects

Contact

Client Portal

Schedule Consultation

ENDPOINT: /api/v1/inquiries

ENTERPRISE ENGAGEMENT DESK

Let's Build Something Exceptional Together

Schedule a technical consultation with our engineering directors or submit detailed project specifications for rigorous architecture review.

Direct Line

15 min avg

Chat on WhatsApp: +1 415 555 0192

TELEMETRY /API/V1/HEALTH

99.98% SLA

Rolling 30-Day Metric

0 Incidents Logged

30 days ago

Operational (Latency 14ms)

Client SLA Guarantee

Mutual NDA signed within **24 business hours**. All briefs receive private architectural triage led directly by a Principal Systems Architect.

DIRECT CHANNELS

Core Engineering

engineering@olinethra.com

Strategic Ventures

partnerships@olinethra.com

GLOBAL ENGINEERING HUBS

3 Worldwide Nodes

San Francisco (HQ)

500 Howard St, Suite 400

PST (UTC-8)

London Hub

100 Bishopsgate, Level 19

GMT (UTC+0)

Singapore Node

Marina Bay Financial Tower 2

SGT (UTC+8)

Decentralized 24/7 Deployment Coverage

Production Project Brief

POST /api/v1/inquiries

Fill out your specifications. Our pipeline ingests your requirements, generates an initial architectural analysis, and routes to the relevant practice lead.

Full Name *

Work Email *

Dr. Elena Vance

e.vance@blackmesa.ai

Company / Organization *

Current Infrastructure / Website URL

Synthetix Dynamics

https://synthetix.io

Primary Engineering Scope *

Select an architectural vertical...

Capital Allocation (USD Budget Range) *

\$50k - \$100k Recommended

\$25k - \$50k

\$50k - \$100k

\$100k+

MVP / Sprint 0

Core System

Enterprise Tier

☒ Under 4 Weeks

☐ 1 - 3 Months

☐ Ongoing Strategic

Technical Specifications & Constraints *

Markdown supported

Detail your technical challenge, current stack (e.g. Go, Rust, PyTorch, Kubernetes), throughput targets, and performance bottlenecks...

Architecture Diagrams / RFPs (Optional)

Attach RFP, PDF, Figma, or System Diagram

Max 45MB. Encrypted in-flight via AES-256

Submit Technical Brief for Architect Review

Zero IP Transfer

Direct Encrypted Pipeline

Senior Staff Architect Triage

Zero-Obligation Deep Dive

We do not employ non-technical account executives. Your consultation is conducted directly with a seasoned software architect who has delivered multi-million request/sec infrastructure.

Live Autonomous Pre-Brief Agent

v1.4


Query our live assistant regarding team capacity, tech stack compatibilities, and engagement rates before booking.

Context Memory: Olinethra-Core-2025

Olinethra AI Agent • 0.02s

Greetings. I can answer technical queries regarding our current engineering bench, distributed consensus models, Kubernetes multi-region failovers, or schedule direct architect consultations. What is your system's primary challenge?

Page 13 of 21



OLINETHRA

High-performance distributed systems, low-latency microservices, and specialized autonomous AI agents engineered for enterprise-grade workloads.

Core API Gateway: 99.99% Uptime

SOLUTIONS

Autonomous Agents
Distributed Pipelines
Neural Search Engines
Edge AI Deployment
Security Auditing

ENGINEERING BRIEFING

Receive our bi-weekly architectural dispatches covering production AI workloads, memory optimizations, and system design.

engineer@domain.com

Subscribe

Design system

The look is light, clean and technical: soft blue-white surfaces, one strong blue, and orange only for calls to action. All tokens live in the `tailwind.config` block at the top of every page and follow Material 3 naming.

Colors

Token	Hex	Use it for
primary	#004fcb	Links, active nav, icons, main buttons
primary-container	#0265ff	Hover states, chat button, highlights
secondary-container	#fe6a17	Orange call-to-action buttons and live status dots only
tertiary	#006178	Small technical labels and eyebrow text
background / surface	#faf8ff	Page background
surface-container-lowest	#ffffff	Cards
surface-container-low	#f2f3ff	Alternate section bands, inner panels
surface-container-high	#e5e7f6	Darker bands, progress tracks

Token	Hex	Use it for
<code>on-surface</code>	#171b26	Headings and main text
<code>on-surface-variant</code>	#424656	Body and secondary text
<code>outline-variant</code>	#c2c6d8	Dividers and separators

Rule of thumb: blue means "information", orange means "act now". Never use orange for decoration.

Typography

Style	Font	Size / line height	Use
<code>display-hero</code>	Plus Jakarta Sans 800	64 / 72px (40 / 48 mobile)	Home hero only
<code>headline-xl</code>	Plus Jakarta Sans 700	48 / 56px	Page titles
<code>headline-lg</code>	Plus Jakarta Sans 700	32 / 40px	Section titles
<code>headline-md</code>	Plus Jakarta Sans 600	24 / 32px	Card titles
<code>headline-sm</code>	Plus Jakarta Sans 600	20 / 28px	Small titles, logo
<code>body-lg</code>	Inter 400	18 / 28px	Hero subtext
<code>body-md</code>	Inter 400	16 / 24px	Default text
<code>body-sm</code>	Inter 400	14 / 20px	Card text, lists
<code>label-md</code>	Inter 600	14 / 20px	Buttons, nav
<code>code-sm</code>	Inter 500	13 / 18px	Endpoint labels, metrics, badges

Spacing and shape

- Spacing scale: `space-xs` 4px, `space-sm` 8px, `space-md` 16px, `space-lg` 24px, `space-xl` 40px, `space-2xl` 64px.
- Page side margin: 20px mobile, 32px tablet, 48px desktop. Content max width is `max-w-7xl` (1280px).

- Corners: 4px default, 8px for buttons and inputs, 12px for cards, full round for pills.
- No visible borders. Separation comes from background shifts and soft shadows (`shadow-sm` , hover `shadow-md`).

Icons and recurring components

- Icons: Google Material Symbols Outlined (`terminal` , `verified` , `hub` , `psychology` , `shield` , `bolt`).
- **Header:** fixed, 80px tall, frosted glass (`bg-surface/80 backdrop-blur-xl`), logo + "API Operational" pill + nav + orange "Schedule Consultation" button.
- **Eyebrow label:** small uppercase `code-sm` text with an icon or pulsing dot above each section title.
- **Card:** white, 12px corners, 24px padding, icon tile top-left, code-style tag top-right.
- **Metric tile:** big number + small label, often in a strip of four.
- **Tech badge:** small grey pill with a stack name.
- **Floating AI button:** bottom-right, blue pill, on every page.
- **Footer:** logo + tagline, Solutions list, Navigation, newsletter signup.

On a phone (390px wide) the nav links disappear with no menu button to reach them, and the hero headline fills the whole first screen. Both need fixing:



OLINETHRA
ENGINEERING &
AI

Schedule
Consultation



● ENTERPRISE CLOUD ARCHITECTURE & AI
RUNTIME

Engineering Scalable Software & Intelligent Cloud Systems

for Ambitious Brands

We architect next-
platforms, enterprise AI solutions, and



AI Agent /api/v1



Interactive behaviors

All interactions today are front-end simulations in plain JavaScript. Nothing is sent to a server yet; each one needs a real backend before launch.

Page	Feature	What it does now	What it needs for real
Home	Terminal tabs (<code>switchTerminalTab</code>)	Swaps the JSON shown for <code>/health</code> vs <code>/services</code>	Optional: real status data
Home	AI chat drawer (<code>toggleAIChat</code> , <code>sendQuickPrompt</code>)	Opens a panel; quick prompts give canned replies	An LLM-backed chat endpoint
Services	Pillar filter tabs (<code>filterPillars</code>)	Shows one of the four pillar cards, or all	Nothing — complete
Services	Scope estimator (<code>calculateEstimates</code>)	Updates sprints, team pod, launch sprint, SLA and security label	Agree real pricing rules; pass the chosen scope into the contact form

Page	Feature	What it does now	What it needs for real
Projects	Category pills + search (<code>applyFilters</code>)	Filters six cards; shows "N shown" and an empty state	Load projects from a CMS or JSON file
Projects	Architecture modal	Opens a detail panel filled from the card	Real write-ups per project
Contact	Budget selector (<code>selectBudget</code>)	Highlights a tier and updates the hint text	Nothing — complete
Contact	Project brief form (<code>handleInquirySubmission</code>)	Shows a spinner, then a fake "#INQ-9482" success message	Send to email / CRM, validate, store uploads
Contact	File upload	Shows the chosen file name	Real upload with size and type checks
Contact	Pre-brief AI agent (<code>handleChatInput</code>)	Echoes the question with a generic reply	Same chat endpoint as Home
All	Floating "AI Agent" button	Visual only on Services and Contact	Wire to the shared chat drawer

Content and voice

Write like a senior engineer talking to another engineer: confident, specific, short. When in doubt, swap an adjective for a number.

- **Do:** "<60s automated recovery", "NDA signed within 24 business hours", "Talk directly to an architect."
- **Don't:** "world-class", "cutting-edge", "synergy", or claims with no figure behind them.
- Use endpoint-style eyebrow labels (`GET /api/v1/...`) for section headers, but keep the headline itself in plain English.
- Every section ends with a clear next step.

Placeholder content to replace before launch

All names, numbers and clients in the mockups are sample content. Nothing should go live until it is replaced with real, verifiable information.

- **Case studies and clients:** FinScale, Aether, Nexus / OmniRoute, Synthetix, PulseVibe, Novus are invented.
- **Testimonials:** Marcus Reynolds, Elena Vargas and Siddharth Kumar are invented people. Use real quotes with permission only.
- **Metrics:** 99.999% uptime, \$1.8B+ processed, 148 microservices, SOC-2 / HIPAA claims. Only publish figures and certifications we actually hold.
- **Contact details:** phone +1 415 555 0192, the San Francisco / London / Singapore addresses, and the two email inboxes.
- **Form demo values:** "Dr. Elena Vance" placeholder and the #INQ-9482 success ID.
- **Copyright:** footer says © 2025; update to the current year.

Intern tasks and next steps

The mockups are static HTML with Tailwind from a CDN. The goal is to rebuild them as a Next.js app in this repo (`portfolio_Nalini`) with shared components. Read `AGENTS.md` first: the repo uses a newer Next.js with breaking changes, so check `node_modules/next/dist/docs/` before writing code.

Fix first

- ☐ Merge `index.html` and `code.html` — they are the same page (only quote marks differ).
- ☐ Make the Services page reachable. Today every nav link for "Services" goes to Home; nothing links to `code3.html` .
- ☐ Make the nav the same on every page. Services shows 7 links (Team, Insights, Careers, FAQs included); the others show 3. Team, Insights, Careers and FAQs have no pages yet.
- ☐ Add a mobile menu button, and use the smaller `display-hero-mobile` size for the hero headline on phones.
- ☐ Point "Client Portal", "Privacy", "Service SLA" and "API Telemetry" at real pages or remove them.

Build in Next.js

1. Move the Tailwind tokens into the project's Tailwind config so every page shares them.

2. Build shared components: Header, Footer, AI chat button and drawer, Card, Metric tile, Tech badge, Eyebrow label.
3. Build the four pages as routes: / , /services , /projects , /contact .
4. Move case studies, services and testimonials into data files so cards render from one source.
5. Rebuild each interaction from the table above as a client component.
6. Connect the contact form to a real endpoint and pass the estimator's scope into it.
7. Check mobile layouts at 375px and keyboard and screen reader access.

Open questions for the founder

- Which case studies, metrics and certifications are real and approved to publish?
- Should the AI chat be a real assistant at launch, or hidden until ready?
- Are Team, Insights, Careers and FAQs in scope for version 1?